

# Modality-Driven Search with Holistic Trace Judging for ARC-AGI-2

Johan Land

2026-02-05

## Abstract

Large language models (LLMs) can generate convincing reasoning traces for abstract reasoning tasks while still being confidently wrong. The core challenge is therefore not only generating solutions, but reliably distinguishing correct from incorrect solutions in the presence of persuasive overclaiming, including within detailed traces. I introduce a solver for ARC-AGI-2 built around two principles: (i) modality-driven candidate generation across independent reasoning channels (text, images, and code with tool-use) to produce diverse candidate solutions, and (ii) context-preserving holistic judging that evaluates all candidate traces jointly in one prompt to identify the most likely correct outputs, including the common case on hard tasks where the correct solution is not the majority hypothesis.

On the ARC Prize semi-private evaluation set (used for the official leaderboard), this solver achieves 72.9% solved at \$38.99/task. On the public evaluation set, it achieves 76.11% solved at \$19.69/task (self-measured). I release the full source code and document negative results showing which common prompting and decomposition strategies reduced diversity and hurt performance.

## 1 Introduction

A central challenge in applying LLMs to abstract reasoning is not just producing candidate solutions, but **knowing what is right and what is wrong** in a setting where models can be confidently incorrect—even when they provide detailed, plausible reasoning traces.

ARC-AGI-2 was designed to be *easy for humans and hard for AI*, and—critically—to measure both **capability** and **efficiency** (cost). Progress on ARC-style benchmarks has also been rapid rather than stagnant: ARC Prize reports significant year-over-year improvements driven by frontier “reasoning systems” and application-layer refinement harnesses (Chollet et al. 2024).

This paper describes an approach that treats **modalities as search operators** and uses **judging as the final selection mechanism**: generate diverse candidate solutions across independent reasoning channels, then select among them using context-preserving holistic judging.

### 1.1 Contributions

- **A modality-driven search solver** that generates candidates independently across text, image, and code reasoning channels to produce diverse candidate solutions.
- **A context-preserving holistic judge** that reads all candidate traces jointly to select the best outputs. Unlike standard self-consistency (majority vote) or per-candidate scoring, this judge identifies correct *minority* hypotheses by comparing full reasoning traces in a single context window — yielding +7 solved instances over majority vote at only 13% of total system cost (Section 7).

- **Verified ARC-AGI-2 semi-private performance:** 72.9% at \$38.99/task on the ARC Prize Verified leaderboard dataset.<sup>1</sup>
  - **Public eval performance:** 76.11% at \$19.69/task (self-measured).
  - **Fully AI-generated implementation:** the entire solver codebase was generated by AI coding assistants (Codex CLI and Gemini CLI); the author did not write or read any solver code. Design ideas were developed through structured debate between ChatGPT-5.2-Pro and Gemini 3 Deep Think, then implemented by the coding assistants (Section 3.4).
  - **Open-source release** of the full source code<sup>2</sup> plus detailed negative results documenting what did not work and why. The complete public-evaluation run data — over 7 million lines of prompts, responses, reasoning traces, and judge transcripts — is also released<sup>3</sup> to support future research.
- 

## 2 Background and Related Work

This paper sits at the intersection of (i) **abstraction-centric few-shot generalization** as instantiated by ARC-style tasks, (ii) **search-based and neuro-symbolic solvers** that treat ARC as a latent-program induction problem, and (iii) **test-time compute scaling** strategies for frontier LLMs—especially approaches that generate multiple candidate trajectories and then **select, verify, or judge** among them. This section reviews the ARC/ARC-AGI benchmark lineage, highlights major families of solver approaches, and situates “modality search + trace-preserving holistic judging” relative to the most relevant prior work.

### 2.1 ARC and ARC-AGI as a benchmark for abstraction under minimal priors

The Abstraction and Reasoning Corpus (ARC) was introduced by Chollet (2019) as part of a broader argument for measuring intelligence as **skill-acquisition efficiency**—how effectively a system can acquire and apply new skills under constrained experience and priors. ARC’s design emphasizes rapid generalization from a small number of examples, with tasks intended to require only relatively elementary “core knowledge” and to discourage reliance on domain knowledge or internet-scale memorization.

ARC tasks are framed as **few-shot input–output induction**: given a handful of training demonstrations (pairs of grids), the solver must infer an underlying transformation rule and apply it to a held-out test input. The hallmark difficulty is **underspecification**: multiple hypotheses can explain the training pairs, but only a subset will transfer to the test instance, so solvers must cope with an intense “many consistent hypotheses” regime where superficial fit is not enough.

As a concrete example, Figure 1 shows all three training pairs and the test input from task 3dc255db.<sup>4</sup> A human might interpret the shapes as “spaceships”: colored particles sit inside each ship on the exhaust side, and the transformation removes them from the interior and places them on the nose, extending the ship in its direction of travel. The solver must infer this rule — identifying containment, directionality, and the interior/exterior distinction — from only three training demonstrations, then apply it to the unseen test input (bottom row). This task remains

<sup>1</sup><https://arcprize.org/leaderboard>

<sup>2</sup><https://github.com/beetree/ARC-AGI>

<sup>3</sup><https://www.kaggle.com/code/johanland/johan-land-solver-v7-public/comments?scriptVersionId=290052212>

<sup>4</sup><https://arcprize.org/play?task=3dc255db>

unsolved by the solver described in this paper: all 29 candidates failed, and none of GPT-5.2, Gemini 3, or Opus 4.5 produced a correct output.

A recurring theme in ARC research is that the benchmark stresses **compositional abstraction** (e.g., combining multiple latent concepts) and **distribution shift within each task** (training vs. test), rather than data-driven interpolation across a large IID dataset. This is part of what makes ARC resistant to straightforward deep learning approaches trained on the released tasks alone, and it motivates solver families that include explicit search, symbolic representations, or test-time adaptation.

## 2.2 The ARC Prize ecosystem and the evolution to ARC-AGI-2

### 2.2.1 ARC competitions and the “slow progress” era

After ARC’s release, the first major public competition was the Kaggle “Abstraction and Reasoning Challenge” (2020). The best-performing solutions in that era were largely **program-synthesis / DSL-search systems**, and performance improved only gradually for several years—an arc that ARC Prize reports explicitly document when motivating why new benchmark design was needed.

The ARC Prize effort expanded this ecosystem with additional competitive events and a more formalized reporting and verification posture, including an explicit policy around “ARC Prize Verified” scores to reduce confusion arising from incomparable, self-reported results.

### 2.2.2 ARC-AGI-2: design goals, splits, and evaluation protocol

ARC-AGI-2 was introduced as a second-generation benchmark intended to provide a more informative signal at the frontier of reasoning systems. The technical report (Chollet et al. 2024) highlights several goals: maintain the original ARC principles and format, reduce susceptibility to brute-force program search, incorporate **first-party human testing**, and increase “signal bandwidth” (a wider useful range of scores to track progress).

ARC-AGI-2 also formalizes dataset splits and calibration:

- **Training set:** 1000 public tasks spanning a wide range of difficulties.
- **Public evaluation set:** 120 calibrated public tasks.
- **Semi-private evaluation set:** 120 calibrated non-public tasks (used for the live leaderboard and ARC Prize leaderboard).
- **Private evaluation set:** 120 calibrated non-public tasks (used for final contest ranking).

A key protocol detail is the use of **pass@2** scoring, acknowledging that some tasks can contain genuine ambiguity; ARC Prize materials emphasize that the benchmark’s human calibration also used the same “two attempts” framing (e.g., tasks solved pass@2 by at least two humans).

Finally, ARC Prize’s public-facing evaluation culture emphasizes not just raw accuracy but also **efficiency** and comparability (including leaderboard reporting and verification norms).

## 2.3 Classical ARC solvers: DSL program synthesis and enumerative search

The most historically influential “classical” ARC solver family treats tasks as **latent programs** composed from a hand-designed library of primitives (a DSL). The Kaggle 2020 top solutions are widely recognized as belonging to this category: they relied on enumerating candidate transformation chains (sometimes aggressively optimized in low-level languages) to find programs consistent with the demonstrations, and then applying the discovered program to the test input.

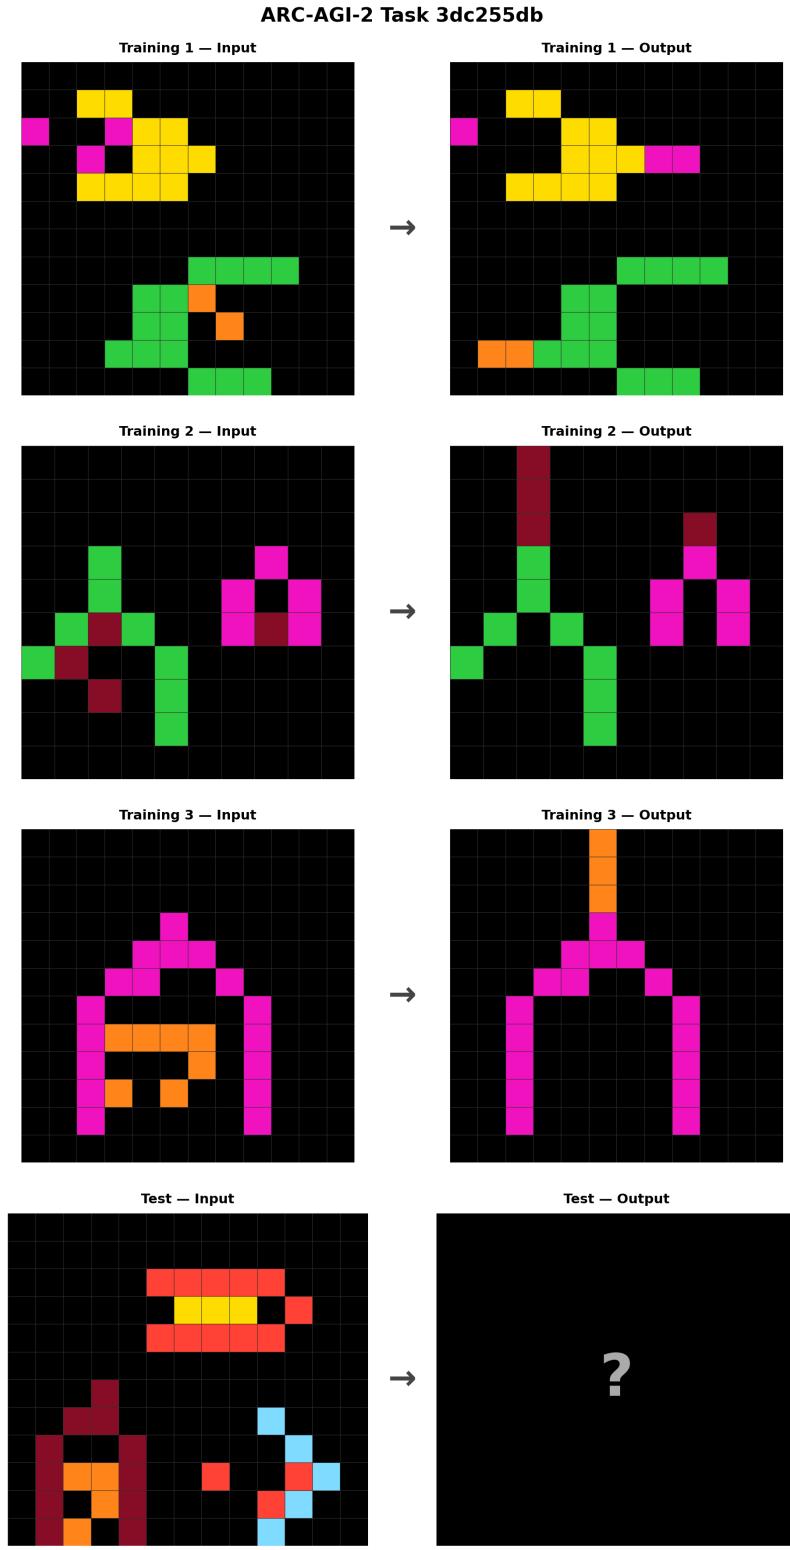


Figure 1: ARC-AGI-2 task 3dc255db. A human might see “spaceships” with particles on the exhaust side. The transformation removes the particles from the interior and extends them from the nose. Three training pairs (rows 1–3) demonstrate the rule; the test input (row 4) must be solved from these examples alone. This task remains unsolved.

These approaches matter as baseline “proofs of tractability” for a subset of ARC tasks and as a reminder that **search can compensate for weak learned priors**—but they also illuminate why ARC-AGI-2 explicitly tries to be “less brute-force.” In particular, ARC Prize’s ARC-AGI-2 announcement and technical report describe removing or redesigning evaluation tasks that were overly susceptible to brute-force search methods.

Beyond DSL enumeration, related symbolic traditions include approaches that use **compression / minimum description length (MDL)** principles to guide search over explanations. Ferré (2021) provides a representative example: describing grids using explicit models and searching for explanations that compress the observations, illustrating an alternative axis of “prior + search” design that emphasizes interpretability and parsimony rather than only brute-force enumeration.

## 2.4 Benchmark extensions and data augmentation in the ARC domain

Several efforts address limitations of the original ARC release: ConceptARC (Moskvichev, Odonard, and Mitchell 2023) introduces concept-grouped task variants to probe systematic generalization; ARC-GEN (Moffitt 2025) proposes a procedural benchmark generator; and Re-ARC (Hodel 2024) provides a programmatic reproduction of ARC tasks frequently used for generating synthetic training variations.

## 2.5 Learned approaches: transduction, induction, and hybridization

ARC has also been attacked from a purely learned perspective (e.g., treating ARC as image-to-image translation), but results historically lagged behind symbolic-search systems and humans. ARC Prize’s technical reporting explicitly notes early deep-learning baselines performing very poorly on ARC-AGI, motivating hybrid approaches and new data strategies.

Recent work has sharpened the conceptual distinction between:

- **Induction:** infer a latent function/program consistent with demonstrations, then apply it.
- **Transduction:** directly predict the test output conditioned on the demonstrations and test input, without explicitly representing a latent program.

Li et al. (2024) study this tradeoff on ARC and find induction and transduction succeed on different problem families; importantly, they show that combining (ensembling) these paradigms can approach human-level performance on the original ARC benchmark under their experimental setup and synthetic training regime. Complementary work on small or specialized models (Fletcher-Hill 2024; Puget 2024) explores architectural inductive biases for 2D grid structure — highlighting the increasing importance of **test-time procedures** (refinement, adaptation, search) even when the base model is learned.

## 2.6 Test-time adaptation and compute scaling for ARC-style tasks

A major recent shift in ARC solving is the move from purely “static” models or solvers to methods that treat each ARC task as an opportunity for **test-time learning** or **test-time search**.

Akyürek et al. (2024) demonstrate that updating model parameters at test time (under carefully controlled procedures) can yield surprisingly strong gains on ARC-like reasoning, reinforcing the idea that ARC is a testbed for *within-task* adaptation rather than only pretraining.

Other ARC Prize-era work explores search and induction at test time in more explicitly programmatic spaces: Macfarlane and Bonnet (2024) combines learned representations with explicit search over programs, while Ouellette (2024) compares neurally-guided program induction

paradigms across grid, program, and transform spaces. At the systems level, competitive solvers increasingly resemble **pipelines** that integrate synthetic data, model adaptation, search components, and ensembles (Chollet et al. 2024).

## 2.7 LLM-era reasoning: generating and coordinating multiple trajectories

Independently of ARC, the broader LLM literature has developed a family of **test-time compute** and **trajectory diversification** techniques that are directly relevant to the solver architecture in this paper. Snell et al. (2024) show that optimally allocating test-time compute — e.g., by generating and selecting among multiple candidate solutions — can be more effective than scaling model parameters, providing formal grounding for architectures that trade inference-time search budget for performance.

### 2.7.1 Chain-of-thought and sampling-based diversification

Chain-of-thought prompting (Wei et al. 2022) established that eliciting intermediate reasoning steps can improve performance on multi-step tasks.

Self-consistency (Wang et al. 2023) then proposed a simple but influential extension: sample multiple reasoning paths and select the most consistent final answer, demonstrating that *diversity + aggregation* can outperform a single greedy reasoning trace.

### 2.7.2 Reasoning as search over “thoughts”

Tree of Thoughts [ToT; Yao, Yu, et al. (2023)] reframes inference as explicit search over intermediate “thoughts,” enabling branching, lookahead, and backtracking. Graph of Thoughts [GoT; Besta et al. (2024)] generalizes this idea to arbitrary graph-structured reasoning artifacts, emphasizing more flexible dependency structures across intermediate units of information.

These frameworks provide language for understanding ARC solvers that “branch” over hypotheses rather than commit early, and they motivate treating candidate generation as a search process rather than a single pass.

### 2.7.3 Tool use and program-aided reasoning

ReAct (Yao, Zhao, et al. 2023) interleaves reasoning traces with actions (tool calls / environment interactions), highlighting how external tools can mitigate hallucination and enable more reliable task completion.

Toolformer (Schick et al. 2023) provides a training-time perspective, showing that LMs can learn to decide *when and how* to call tools via self-supervision, further legitimizing tool-augmented reasoning as a general capability axis.

PAL [Program-aided Language Models; Gao et al. (2023)] formalizes a closely related idea: use the LM to translate problems into runnable code, then outsource exact computation to an interpreter—often improving accuracy on tasks where arithmetic/logic errors dominate.

These ideas are directly relevant to ARC, where code synthesis can serve both as a hypothesis generator and as a way to expose structured intermediate artifacts (programs, tool outputs) that may be useful during downstream selection.

### 2.7.4 Iterative refinement and memory-based improvement

Self-Refine (Madaan et al. 2023) shows that a single LM can iteratively improve its output by generating feedback and revisions in a loop, without gradient updates. Reflexion (Shinn et al.

2023) similarly aims to improve test-time performance by incorporating feedback into a memory buffer that influences subsequent attempts, framing improvement as “verbal reinforcement learning” rather than weight updates.

These methods relate to ARC attempts that “try multiple times” and learn from earlier mistakes, though they also introduce a key trade-off: iterative refinement can increase compute while risking **anchoring** to early hypotheses—an issue that becomes acute on ARC tasks where early commitments can be misleading.

## 2.8 Selection, verification, and “LLM-as-a-judge” paradigms

Generating diverse candidates is only half the problem; the other half is selecting among them. In the LLM ecosystem, selection is increasingly delegated to learned or LLM-based evaluators, giving rise to the “LLM-as-a-judge” paradigm.

Zheng et al. (2023) formalized and stress-tested LLM judging in the context of MT-Bench and Chatbot Arena, showing that strong LLM judges can correlate well with human preferences while also exhibiting systematic biases (e.g., position and verbosity biases). Subsequent work has explored structured multi-agent evaluation (judge-and-jury designs) and cautioned that judging format (pairwise vs. pointwise, aggregation schemes) materially affects robustness.

For ARC in particular, selection is unusually difficult because many hypotheses fit the demonstrations yet fail on the test instance. A good judge must reward **transfer-valid abstractions**, not merely fluent rationales or training-pair fit.

## 2.9 Positioning of this work within the landscape

Relative to the above literature, the approach in this paper can be viewed as a systems-level composition of three trends:

1. **Hypothesis generation as explicit search** (a lineage shared with DSL/program-synthesis solvers and ToT/GoT-style inference (Yao, Yu, et al. 2023; Besta et al. 2024)), but broadened beyond a single representation to *multiple reasoning modalities*.
2. **Tool- and program-mediated reasoning** (Yao, Zhao, et al. 2023; Gao et al. 2023; Schick et al. 2023) used not just for execution but also as a way to produce richer intermediate artifacts that can be consumed by downstream selection.
3. **Judge-based selection** [LLM-as-a-judge; Zheng et al. (2023)] adapted from general LLM evaluation into an in-task meta-reasoning component, with the additional complication that ARC demands judging *generalization under underspecification*, not merely surface quality.

What is relatively distinctive in the ARC context is the combination of (i) **heterogeneous candidate generators** (text, code, visual, extended deliberation) and (ii) **context-preserving comparison over full traces**, rather than scalar scoring or consensus compression—an axis that is motivated both by historical ARC solver failure modes (overfitting via brittle heuristics) and by known limitations of LLM judges under compressed, bias-prone evaluation formats.

### 3 Method: Modality-Driven Search and Architecture

#### 3.1 Core idea: independent candidates across modalities maximize diversity

The solver is built around a practical observation: to solve tasks that frontier systems and labs *do not already solve*, the correct solution is often a **minority hypothesis**. If “the most common solution” were correct, the task would likely already be within the main cluster of model behavior.

Therefore, the solver’s first phase intentionally creates many *independent* candidate solutions across heterogeneous reasoning modalities (text, image, and code), maximizing the probability that at least one candidate captures a genuinely novel hypothesis. Each modality provides a structurally different representation of the same task, which empirically produces candidates that cluster differently (Section 6.5).

#### 3.2 Pipeline overview

Figure 2 shows the end-to-end pipeline. For each ARC task:

1. **Candidate generation (up to 29 candidates).** Run a set of modality-specific solvers, each producing:
  - a single predicted output grid,
  - a reasoning trace,
  - for codegen-with-tools candidates: the code, tool calls, and tool outputs (execution logs).

Each candidate produces one output; the pass@2 two-guess format is introduced at the judging stage (step 3). Candidate generation proceeds in stages: if early-stage candidates already show strong agreement (multiple candidates converging on the same output), the system terminates early and skips the remaining, more expensive modalities. This adaptive early stopping improves cost efficiency on tasks that can be solved with a shallower search. On ARC-AGI-1 (which contains easier tasks that more frequently trigger early stopping), the same system achieves 94.5% on the official semi-private evaluation at only \$11.40/task — substantially cheaper than the \$38.99/task on ARC-AGI-2, where harder tasks require the full candidate budget more often. Both semi-private scores and costs are as reported by ARC Prize’s verification infrastructure.<sup>5</sup>

2. **Holistic judging (3 parallel judges).**

Concatenate all candidate traces into a single long-context prompt and ask a judge model to:

- identify the top-2 most likely correct candidates (or propose a synthesis),
  - explain why other clusters are wrong,
  - output the final grids.
3. **Weighted scoring.** Each judge’s first choice receives 2 points and second choice receives 1 point. The two distinct output grids with the highest total score become the solver’s pass@2 guesses.

---

<sup>5</sup>The semi-private evaluation is run by ARC Prize on non-public tasks; the author does not control that environment. Because the semi-private logs are not available for inspection, the detailed analysis in this paper (Sections 6–8) is based on the public evaluation run, where full data is available.

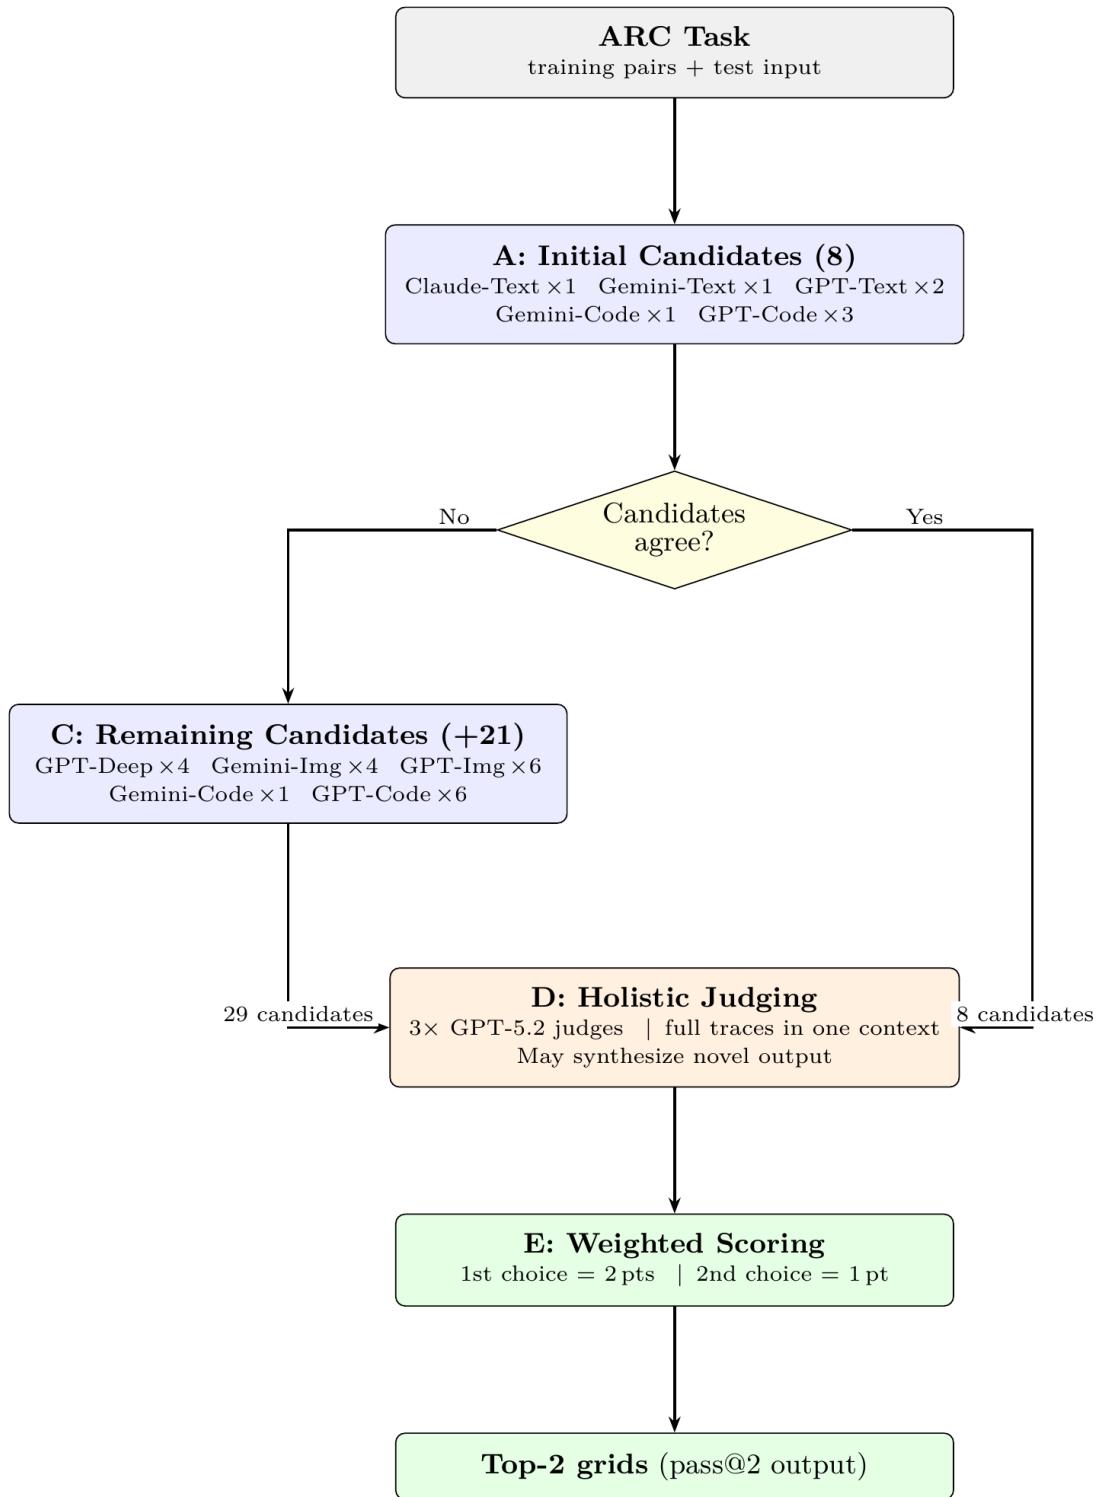


Figure 2: Solver pipeline: candidate generation with adaptive early stopping, holistic judging, and weighted scoring.

### 3.3 Development methodology: AI-generated design and implementation

The solver described in this paper was built entirely through AI-assisted development. The author’s role was strategic direction — deciding what to build, evaluating results, and selecting next steps — but **no solver code was written or read by the author**. All implementation was performed by AI coding assistants (Codex CLI and Gemini CLI), operating from natural-language specifications.

The design ideas themselves originated from a structured process: the author posed architectural and strategic questions to ChatGPT-5.2-Pro and Gemini 3 Deep Think, often having the two models debate each other on topics such as candidate diversity strategies, judge prompt design, and failure mode analysis. The resulting ideas were refined through this adversarial dialogue and then handed to the coding assistants for implementation.

This development pattern — AI models generating the ideas, AI coding assistants implementing them, with a human orchestrating the process — represents a recursive application of the same “generate diverse candidates, then select” principle that underlies the solver itself. It also suggests that frontier-level systems engineering may increasingly be achievable through AI-mediated design, even for novel and complex pipelines.

### 3.4 Why candidate diversity matters specifically on ARC-AGI-2

ARC-AGI-2 tasks often introduce new concepts. A solution can be perfectly “logical” on the training pairs but still be a brittle overfit that fails to abstract the intended rule. This makes naive “logic checking” of traces insufficient (details in Section 5 and Section 8). In practice, the system needs:

- broad hypothesis exploration across structurally different representations, and
  - a judge that can identify **where** a candidate is likely overfitting—even when it reads as coherent.
- 

## 4 Method: Multimodal Candidate Generation

### 4.1 Models and inference configurations used

Candidate generation is performed via three foundation models and multiple inference configurations:

- **Gemini 3 Preview**, high-reasoning setting (text and image; codegen with tools).
- **GPT-5.2**, x-high reasoning setting (text; codegen with tools and without tools; image; “deep thinking” configuration).
- **Claude Opus 4.5**, long-context (120k) (text reasoning).

*(Exact API parameters, tool schemas, and prompts are documented in the open-source implementation.)*

In the final system, candidate generators are grouped into three families—**Text**, **Image**, and **Code**—with multiple configurations within each family to encourage diversity. Table 0 shows the full candidate configuration.

**Table 0. Candidate configuration: 29 generators grouped by family.**

Family	Generator	Candidates
Text	Claude Opus 4.5 (text)	1
Text	Gemini 3 Preview (text)	1
Text	GPT-5.2 (text)	2
Text	GPT-5.2 (deep think)	4
Image	Gemini 3 Preview (image)	4
Image	GPT-5.2 (image)	6
Code	Gemini 3 Preview (code, tools)	2
Code	GPT-5.2 (code, tools)	9
<b>Total</b>		<b>29</b>

The text family contributes 8 candidates (including 4 deep-think runs), image contributes 10, and code contributes 11. Within each family, multiple runs of the same generator use the same prompt and API parameters; diversity arises from model sampling stochasticity.

## 4.2 Text methodologies

Text candidates are generated by prompting a language model with a textual encoding of the ARC training pairs and test input. The model is asked to infer the transformation rule and to output a completed test grid. Text prompting is used with multiple foundation models (Gemini/GPT/Opus) as independent candidates.

### 4.2.1 Base text prompt

The base prompt is intentionally minimal:

```
You are solving an ARC (Abstraction and Reasoning Corpus)
task. Each grid cell is an integer 0-9 representing a color.
Use the solved examples to infer the transformation and
apply it to the test input.

...
{training and test examples}
...
Respond with an explanation of your thinking that is detailed
enough that someone can reconstruct your solution. Afterwards,
you MUST also respond with the completed output grid.
```

The prompt deliberately does **not** prescribe a fixed reasoning template, a step-by-step plan, or a fixed output grid format. This reduces “prompt compliance” overhead and empirically increases hypothesis diversity. The trade-off is that outputs are noisier and require tolerant parsing and validation to recover candidate grids (see also Section 8 for supporting negative results on strict output constraints and prescribed reasoning templates).

Grids are encoded in **CSV format**, which was selected after benchmarking 9 representation formats (standard space-separated, semicolon-delimited, XML-tagged, CSV, Python lists, sparse coordinate notation, ASCII symbols, binary masks, and compact pipe-delimited). Suboptimal format choices cost on the order of 10% lost performance relative to CSV, with compact formats that are difficult for LLMs to produce (e.g., sparse coordinate notation, binary masks) performing substantially worse.

### 4.2.2 “Deep think” variant

In addition to the standard text prompt, I run a “deep think” configuration that allocates a larger test-time compute budget. In practice, this is achieved via a prompt modification that explicitly encourages GPT-5.2 to reason more extensively before committing to a final output — effectively trading tokens for deeper deliberation within a single response.

## 4.3 Image methodologies

Image candidates are generated by rendering the ARC training pairs and test input as a single annotated image and providing it alongside the (otherwise similar) instruction prompt. This provides the model with a pixel-space representation that can be advantageous for tasks where the salient structure is more readily perceived visually than through a textual grid encoding.

Notably, the image renderings are intentionally **imprecise** — slightly distorted rather than pixel-perfect grid reproductions. Pixel-perfect renderings underperformed in early experiments, possibly because models treated them as lossless encodings and fell back on cell-by-cell numerical reasoning rather than engaging the visual pattern recognition that makes image prompting valuable in the first place. The slight imprecision appears to encourage models to reason about shapes, symmetries, and spatial relationships at a higher level of abstraction.

Figure 3 shows an example rendering for task `d35bdbdc:1` (public evaluation split), where each training pair is shown as an input/output image pair.

In the public evaluation analysis (Section 6.6), image prompting provides uniquely correct candidates on several instances, including: `20a9e565:1`, `2d0172a1:1`, `4e34c42c:1`, `b6f77b65:1`, `b6f77b65:2`, and `d35bdbdc:1`.

## 4.4 Code methodologies

Code candidates treat ARC solving as program synthesis: the model is asked to produce executable code that maps input grids to output grids. I use two code-generation regimes:

1. **Tool-integrated code generation (native tool calls).** The model iteratively writes code, executes it via a sandbox tool, inspects intermediate outputs, and refines the program across multiple tool calls.<sup>6</sup>
2. **One-shot code generation (no tools).** The model returns a complete program in a single response; the harness executes it in a sandbox and records the result, but the model does not receive iterative execution feedback.

The tool-integrated regime often produces rich intermediate artifacts (program drafts, test harnesses, and execution traces) that are later consumed by the holistic judge. Task `13e47133:1` is a representative example where code-based reasoning dominates (Section 6.6). A shortened excerpt of the tool-integrated trace illustrates the iterative development pattern:

```
"detailed_logs": [
  {
    "type": "code",
    "code": "import numpy as np, collections,
```

<sup>6</sup>The ARC Prize semi-private evaluation uses a Zero-Day Reasoning (ZDR) environment that disables OpenAI tool calls. For the semi-private run, tool-integrated code candidates were replaced with one-shot code generation, reducing the iterative refinement available to code candidates on that evaluation.

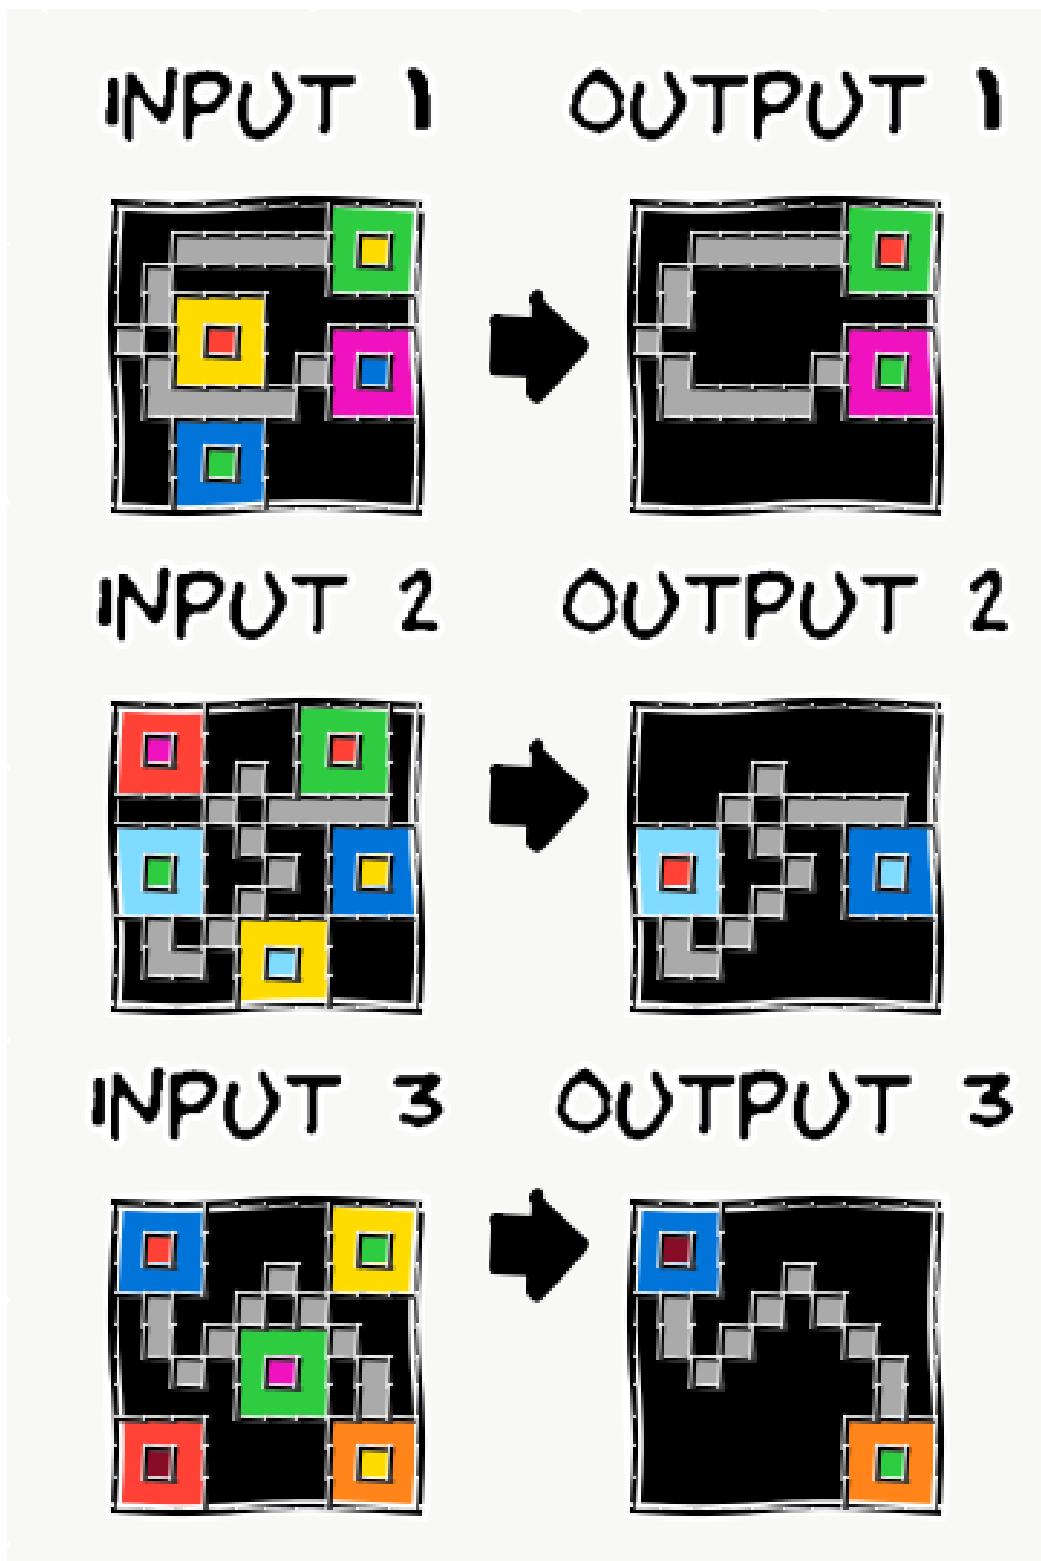


Figure 3: Example image rendering used for image-based prompting (task d35bdbdc:1).

```

        itertools, math, sys, statistics\n",
    "language": "python"
},
...
{
    "type": "code",
    "code": "# Display expected outer right region
# values for rows0-19 cols9-19
outer_expected = ex1_out[:,9:20]
outer_pred = pred1[:,9:20]
outer_in = ex1_in[:,9:20]
outer_expected.shape\n",
    "language": "python"
},
...
{
    "type": "code",
    "code": "# Let's craft final version and re-test.
import numpy as np
from collections import Counter, deque

def solver(input_grid):
    g = np ...\\n",
    "language": "python"
},
{
    "type": "text",
    "content": "### FINAL SOLUTION ###
import numpy as np
from collections import Counter, deque

def solver(input_grid):
    g = np.array(input_grid, ...)
    h, w = g ..."
}
]

```

The one-shot (no-tools) regime is substantially cheaper and can still contribute unique solutions, but it is less robust on tasks that benefit from iterative debugging.

---

## 5 Method: Context-Preserving Holistic Judging

### 5.1 The judging problem: selecting the right needle in a noisy haystack

Given 8–29 candidates, many are plausible and internally coherent. Worse, models tend to **cluster** around the same wrong interpretation on the hardest tasks. The key difficulty is identifying a rare

candidate that is correct—or close enough that it can be repaired—without reducing everything to an overly lossy score.

## 5.2 Judges attempted (and why the “holistic” judge wins)

I tested three approaches:

1. **Logic judge (failed mode):** score candidates by whether their reasoning appears logically consistent.  
Failure mode: a candidate can be “logical” yet overfit to training pairs or miss a newly introduced latent concept in the test case. Therefore logic alone is not robust.
2. **Consistency judge (partial):** look for themes that repeat across candidates.  
Failure mode: consistency tends to reward the majority cluster. But breaking new ground requires elevating *divergent* hypotheses, because the correct solution to an unsolved task is often not the modal answer.
3. **Holistic judge (final):** provide *all traces together* and ask the judge to pick the top-2 most likely correct candidates. Run three judges in parallel and aggregate via weighted scoring (see below). This works because **having all context together beats abstracting traces into scores**. It lets the judge detect subtle but decisive differences between near-identical hypotheses.

The holistic judge can be understood as an **anti-consistency** mechanism. Self-consistency (Wang et al. 2023) selects the most common answer across samples — an effective strategy when the majority is likely correct. On ARC-AGI-2, the majority is often *wrong* on the hardest tasks (Section 6.8), making consistency-based selection a liability. The holistic judge inverts this: it is designed to identify a correct *minority* hypothesis against a confidently wrong majority, using full-trace comparison to distinguish genuine insight from plausible groupthink.

All three judges use **GPT-5.2** (x-high reasoning setting), the same model used for candidate generation but in a distinct role with a different prompt. I also tested a mixed-model judge ensemble (combining Opus, Gemini, and GPT-5.2), but three homogeneous GPT-5.2 judges outperformed the mixed configuration. Using the same model family for both generation and judging introduces a potential correlation risk (the judge may favor candidates “in its own style”), but in practice this was outweighed by GPT-5.2’s stronger individual judging capability.

## 5.3 Judge prompt structure

The holistic judge prompt is assembled programmatically and follows this structure (condensed; the full implementation is in the open-source release):

Below is a problem that was attempted to be solved {N} times:

{training pairs + test input}

Solutions were generated {N} times, using different types of solvers.

```
<SOLUTION 1 START>
<CONTENT>
{full reasoning trace or extracted solver function}
```

```

</CONTENT>
<PREDICTED_GRID>
{candidate output as CSV}
</PREDICTED_GRID>
<SOLUTION 1 STOP>

... (repeated for all N solutions) ...

```

Your task is to understand these solutions, and assess how well they've understood the problem, and how likely their solutions are to provide the correct solution to the test input.

Often, new mechanics are introduced in the test example for which the solutions do not generalize well. Please output two solutions that you think represent the right mechanic for solving the problem.

Output your two solutions as grids (in code blocks). Explain how you came to these two solutions being the two most likely. Study all the provided solutions and their reasoning to come up with a meta-conclusion about how to solve the problem.

Each candidate's CONTENT block contains the full reasoning trace for text/image candidates, or the extracted `solver()` function for code candidates. The PREDICTED\_GRID block contains the candidate's output grid in CSV format. Candidates that produce identical output grids are listed as separate solutions (with separate traces), preserving the judge's ability to assess reasoning quality even when outputs agree. The prompt does **not** instruct the judge to prefer majority or minority answers — it asks for a “meta-conclusion,” leaving the judge free to weigh agreement, reasoning quality, and novelty as it sees fit.

#### 5.4 Allowing synthesis (new solutions not in candidates)

The holistic judge is also permitted to propose a **novel solution not identical to any candidate output**, effectively recombining correct subcomponents of multiple flawed candidates. This matters on tasks where no single candidate “gets it,” but multiple candidates contain partial truths. When synthesizing, the judge outputs a raw output grid directly (not executable code), which means synthesized solutions do not benefit from programmatic verification and are susceptible to arithmetic or grid-construction errors.

#### 5.5 Aggregation: from 3 judges to 2 final guesses

Each judge outputs a ranked pair of solutions (first choice and second choice). To produce the final pass@2 output, the system assigns **2 points** to each judge's first choice and **1 point** to each judge's second choice, then sums points across judges for each distinct output grid. The two grids with the highest total score become the solver's two guesses. When judges agree on their top pick, that solution accumulates up to 6 points; when they disagree, the scoring naturally surfaces the most broadly supported candidates. In practice, full three-way disagreement on the first choice is rare on easier tasks where candidates converge, but becomes common on harder tasks where the judges face the same ambiguity as the generators.

## 5.6 Known limitation: judge biases

Candidate traces are concatenated into the judge prompt in a fixed order; the order is **not shuffled** between judge runs. LLMs exhibit known position biases (e.g., favoring candidates near the start or end of the context), and this ordering could systematically advantage or disadvantage certain candidates. Shuffling the candidate order across the three judge runs and measuring the effect on agreement and accuracy is a natural improvement.

Beyond position bias, the judge may also exhibit **verbosity bias** (favoring longer, more detailed reasoning traces over terse but correct ones) and **format bias** (favoring candidates whose output format more closely matches the judge’s own generation patterns). These biases are well-documented in LLM-as-a-judge settings (Zheng et al. 2023) and could interact with the modality mix, since code candidates tend to produce structured traces while text candidates produce prose. No debiasing is applied in the current implementation.

## 5.7 Feasibility and context length

The holistic judging prompt is intentionally large: on the order of **30k–80k input tokens**, because it includes full traces from many candidates.

This makes long-context frontier models a practical requirement for the judge step.

---

# 6 Experiments and Main Results

## 6.1 Evaluation datasets and protocol

ARC-AGI-2 provides multiple evaluation sets (public, semi-private, private), all calibrated and evaluated under **pass@2** (Chollet et al. 2024).

ARC Prize Verified results are reported on the **semi-private evaluation set** via an official verification process and leaderboard.

**Development data disclosure:** The solver was iteratively designed using both the 1,000-task training set and the 120-task public evaluation set; final configuration tuning (modality mix, candidate counts, judge settings) was performed against the public evaluation split. The semi-private evaluation was run on held-out tasks unseen during development, and the resulting ~3 percentage-point gap (76.11% public vs 72.9% semi-private) suggests limited overfitting to the public split. Additionally, the same system was run on ARC-AGI-1’s semi-private evaluation set (achieving 94.5%) with **no exposure** to ARC-AGI-1 tasks during design — a fully blind evaluation that further validates generalization.

## 6.2 Metrics

I report:

- **Accuracy (pass@2):** the mean per-task solve rate, where each task’s solve rate is the fraction of its test instances answered correctly within two guesses. For tasks with a single test instance this is binary (0 or 1); for multi-instance tasks it can be fractional (e.g., 0.5 if one of two test instances is solved). This is distinct from the *instance-level* accuracy (fraction of individual test instances solved), which is reported separately in the per-instance analyses below.

- **Cost per task (\$/task):** total runtime API cost divided by number of tasks (including candidate generation + judging + tool calls + retries as applicable).

## 6.3 Headline results

### 6.3.1 Timeline

The solver was submitted to the ARC Prize foundation on **December 15, 2025**. Official results were announced on **February 3, 2026**. The leaderboard snapshot in Table 1 reflects the state at the time of announcement; subsequent entries (discussed below) have since been added.

### 6.3.2 Semi-private evaluation (ARC Prize Verified / leaderboard)

My solver achieves:

- **72.9% solved** on ARC-AGI-2 semi-private eval ( $\approx 73\%$ ) at **\$38.99/task**.

For context, the next two leaderboard entries at the time of the results announcement are:

- **GPT-5.2 Pro:** 54.2% at \$15.72/task
- **Gemini 3 Pro:** 54.0% at \$30.57/task

### 6.3.3 Public evaluation (self-run)

On the public eval set, my solver achieves:

- **76.11% solved** at **\$19.69/task** (self-measured).

For the per-instance analysis in Section 6.6, the public evaluation split contains **120 task IDs** with **167 test instances** (75 tasks with 1 test instance, 43 with 2, and 2 with 3).

ARC Prize notes that, in principle, calibrated public/semi-private/private eval sets should be comparable when systems are not overfit (Chollet et al. 2024). The  $\sim 3$  percentage-point gap between public (76.11%) and semi-private (72.9%) likely reflects three factors: (i) natural generalization loss to a held-out task distribution, (ii) the semi-private verification environment used OpenAI’s zero-data-retention (ZDR) API mode, which disables function/tool calling, and (iii) the semi-private run coincided with a period of known instability in OpenAI’s API, resulting in high failure rates and extensive retries that degraded both cost and effective candidate coverage. In this configuration, the tool-integrated code generation candidates (Section 4.4) were replaced with one-shot code generation (no iterative sandbox execution). Code candidates were still produced, but without the iterative debugging loop that makes tool-integrated generation more robust on complex tasks. Since tool-integrated code generation accounts for the bulk of the Code family’s cost (Table 6), the ZDR constraint both reduced accuracy and changed the cost profile of the semi-private run relative to the public run.

**Table 1. Leaderboard snapshot and reference systems.**<sup>7</sup> Semi-private results are as reported on the ARC Prize Verified leaderboard at the time of the official results announcement (February 3, 2026); the public-evaluation row is self-measured on the public evaluation split.

---

<sup>7</sup><https://arcprize.org/leaderboard>

AI System	Author	ARC-AGI-2	Cost/Task	Comment
Human Panel	Human	100.00%	\$17.00	At least two humans out of ~400 solved it
This paper	Johan Land	72.90%	\$38.99	Semi-private (official)
This paper	Johan Land	76.11%	\$19.69	Public eval
GPT-5.2 Pro (High)	OpenAI	54.20%	\$15.72	
Gemini 3 Pro (Refine.)	Poetiq	54.00%	\$30.57	
GPT-5.2 (X-High)	OpenAI	52.90%	\$1.90	
Gemini 3 Deep Think (Preview)	Google	45.10%	\$77.16	
GPT-5.2 (High)	OpenAI	43.30%	\$1.39	
GPT-5.2 Pro (Medium)	OpenAI	38.50%	\$8.99	
Opus 4.5 (Thinking, 64K)	Anthropic	37.60%	\$2.40	
Gemini 3 Flash Preview (High)	Google	33.60%	\$0.23	

In this snapshot, the system described in this paper achieves substantially higher verified semi-private accuracy than the strongest single-entry commercial baselines (72.9% vs. ~54%), indicating that modality-driven candidate generation combined with long-context judging can move the frontier in capability. The remaining gap to the human panel (100.0% at \$17/task) indicates that the benchmark still contains substantial headroom at roughly comparable cost. The accuracy gain over commercial baselines comes at higher cost per task relative to the cheapest entries, reflecting the additional test-time compute spent on multi-candidate search and downstream adjudication. The public-evaluation result (76.11% at \$19.69/task) suggests that comparable accuracy can be achieved at materially lower cost on the public split, although comparisons across public vs. semi-private verification regimes should be interpreted cautiously.

### 6.3.4 Evolving model landscape

The leaderboard is evolving rapidly. Since the official results announcement, Claude Opus 4.6 has been added at 69.2% accuracy and \$3.47/task — approaching this system’s accuracy at a fraction of the cost. For context, the Opus 4.5 model used as a text-reasoning candidate generator in this system scored only 37.6% as a standalone solver; Opus 4.6 represents a near-doubling of single-model performance on this benchmark.

Critically, Opus 4.6 was **not available** during the development of this system (the submission was finalized on December 15, 2025). The ensemble described here was built exclusively with models available at that time. Incorporating stronger base models — such as replacing the Opus 4.5 text-reasoning candidates with Opus 4.6, or using Opus 4.6 as an additional candidate generator — would likely yield a material accuracy improvement over the current 76.11%, since the ensemble’s value scales with the quality and diversity of its constituent generators. This underscores that the

contribution of this paper is the **architecture** (modality-driven search + holistic judging), not the specific model snapshot.

#### 6.4 Efficiency discussion

ARC-AGI-2 explicitly evaluates efficiency; ARC Prize argues cost per task is the most directly comparable efficiency axis across humans and AI systems (Chollet et al. 2024).

Reported cost per task on both evaluation sets:

- **Semi-private (official):** \$38.99/task at 72.9%
- **Public eval (self-measured):** \$19.69/task at 76.11%

For comparison, the next-best entries on the leaderboard at the time of writing:

- GPT-5.2 Pro: **\$15.72/task** at 54.2%
- Gemini 3 Pro: **\$30.57/task** at 54.0%

The roughly 2 $\times$  cost difference between the semi-private and public runs (\$38.99 vs. \$19.69) is likely caused by API-level unreliability during the semi-private run: on the public-eval run, only 2,216 of 14,106 GPT-5.2 API attempts succeeded, with the remainder failing due to rate limits, timeouts, and server errors. The semi-private cost is as reported by ARC Prize’s verification infrastructure and includes costs incurred by failed and retried API calls that do not contribute to the final output. The public-eval cost of **\$19.69/task** is therefore a more representative measure of the system’s actual compute requirements.

At this cost, the solver is comparable to GPT-5.2 Pro (\$15.72/task) while achieving a +21.9 percentage-point accuracy gain (76.11% vs. 54.2%), and is both cheaper and substantially more accurate than Gemini 3 Pro (\$30.57/task at 54.0%). A detailed cost breakdown by component for the public-eval run is provided in Section 7.

#### 6.5 Modality contribution and diversity (qualitative)

The final modality mix was selected based on two criteria:

1. **Performance:** raw solve contribution.
2. **Diversity contribution:** uniquely solved tasks that other modalities fail.

Qualitatively:

- GPT dominates **code generation**, with Gemini adding meaningful diversity; Opus codegen behaved largely like a subset and was dropped in the final mix.
- For **image reasoning**, Gemini and GPT behaved differently and were complementary; Opus image reasoning behaved more like a subset.
- Opus was exceptional for **end-to-end text reasoning**, being the sole solver for several tasks via text-only reasoning.

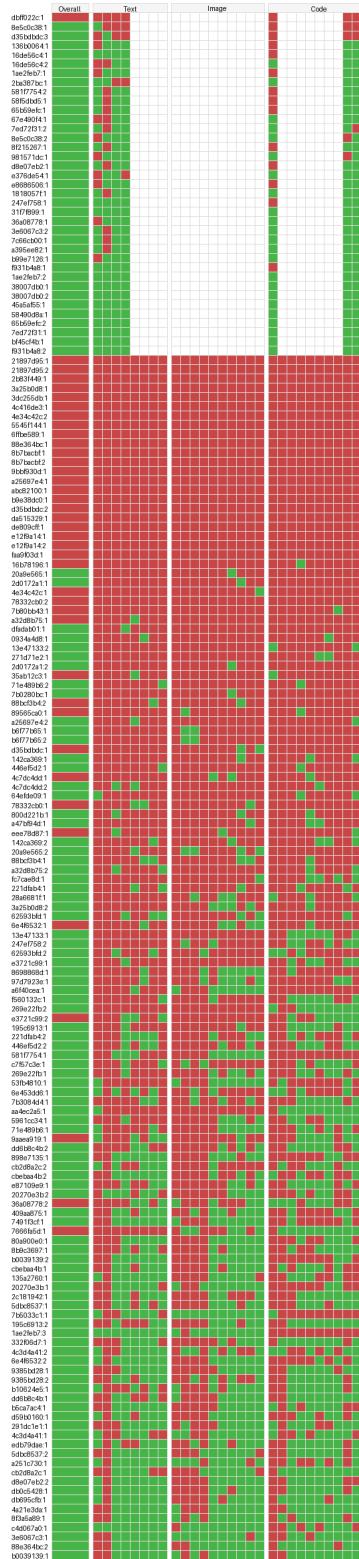


Figure 4: Methodology matrix over public evaluation instances. Green = candidate matches ground truth; red = candidate does not match; white = candidate not produced.

## 6.6 Modality complementarity and uniqueness (public eval)

To quantify complementarity between candidate-generation methodologies, I evaluate each candidate output against the ground-truth test target and record a per-instance correctness matrix (Figure 4). Rows correspond to test instances (`task_id`, `test_index`) and columns correspond to individual candidate generators (model  $\times$  modality  $\times$  configuration). A cell is marked PASS if the candidate output exactly matches the ground truth, and FAIL otherwise. Blank cells indicate that the corresponding candidate generator was not executed for that instance.

Over the 167 public-evaluation test instances, the final system solves **128/167 = 76.65%** at the instance level (pass@2). This is slightly higher than the task-level accuracy of 76.11% reported in Table 1, because partially solved multi-instance tasks pull the task-level average below the raw instance rate. Of 120 tasks, 86 are fully solved (all instances correct), 11 are partially solved (at least one but not all instances correct), and 23 have no correct instances. The candidate pool contains at least one correct output for **144/167 = 86.23%** of instances (candidate-oracle accuracy). The 39 unsolved instances therefore decompose into:

- **22/167 (13.17%)** instances where no candidate in the pool is correct (candidate-generation failures).
- **17/167 (10.18%)** instances where at least one correct candidate exists but is not selected (selection/judging failures).

Notably, there is one instance where the final system output is correct despite **zero** candidates matching the ground truth (21897d95:2); this occurs via judge synthesis (Section 6.7), where the judge recombines partial insights from multiple flawed candidates to produce a novel correct output.

I operationalize modality-level uniqueness by grouping candidate generators into three families: **Text** (including Deep), **Image**, and **Code**. An instance is counted as solvable by a family if any candidate within that family is correct. As described in Section 3, the solver uses adaptive early stopping: when early-stage candidates show strong agreement, the system skips remaining modalities to save cost. This means 37 instances were evaluated with only 8 candidates (the initial Text + Code probe) rather than the full 29. Of these 37 early-stopped instances, **36 were solved correctly** (97.3% accuracy), compared to 92/130 (70.8%) for full-coverage instances. Only one early-stopped instance (`dbff022c:1`, <https://arcprize.org/play?task=dbff022c>) was incorrectly solved — a case of extreme groupthink. The task itself is relatively simple, but the test case introduces a new mechanic with two valid interpretations: the legend that maps symbols to colors is either “the same” as in training or “inverted.” All models confidently assume the simpler interpretation (“legend is the same”), while the ground truth requires the more complex one (“legend is inverted”). This is arguably an artificial source of difficulty — the ambiguity is not clearly disambiguated by the training examples — but it illustrates a characteristic challenge of ARC: test cases can introduce subtle twists that render the majority hypothesis wrong, and no amount of additional candidates would help when every model makes the same simplifying assumption. This high accuracy rate (36/37) validates the early-stopping heuristic, but only for the easiest tasks — which are precisely the tasks that trigger early stopping. On harder tasks, candidates do not converge, early stopping is not triggered, and the full 29-candidate budget is spent. It is on these harder tasks that the system’s core contributions — modality diversity and holistic judging — become essential, because majority voting systematically fails when the correct solution is a minority hypothesis (Section 5, Section 7).

To avoid conflating uniqueness with conditional execution, I restrict the modality-level analysis below to the **130 instances with complete candidate coverage** (29/29 candidate columns filled).

**Table 2. Pairwise non-overlap between modality families on the complete-coverage subset ( $n = 130$ ).** Each entry counts instances with at least one PASS in the row family and zero PASS in the column family.

	Text	Image	Code
Text	NA	13	7
Image	17	NA	11
Code	18	18	NA

**Table 3. Exclusive coverage on the complete-coverage subset ( $n = 130$ ).** An instance is exclusive to a family if that family has at least one PASS and the other two families have zero PASS.

Family	Count
Text only	2
Image only	6
Code only	7

Taken together, Tables 2–3 indicate substantial complementarity between modality families: each family covers a non-trivial set of instances that are not covered by at least one of the other families, and exclusive coverage persists even when all candidates are generated. This supports treating modalities as distinct search operators rather than relying on a single representation.

Beyond these aggregate counts, Figure 4 exhibits pronounced *instance-level heterogeneity*: some tasks are solved reliably by one family while being largely unsolved by others. For example, task `a6f40cea` (<https://arcprize.org/play?task=a6f40cea>) is solved by **7/10** Image candidates, **1/8** Text candidates, and **0/11** Code candidates on its test instance. Qualitatively, the underlying rule resembles a localized magnification or “lens” operation (identifying a small region and “zooming” it), which is naturally expressed and discovered in pixel space. Conversely, task `13e47133` (<https://arcprize.org/play?task=13e47133>) is solved *only* by Code candidates on both of its test instances (Text = 0/8, Image = 0/10, Code = 7/11 and 2/11, respectively). In this case, the solution is readily implemented by explicitly traversing the boundary/edge structure, a representation that is directly available to code-based generators but comparatively indirect for pure text or visual prompt-based reasoning.

## 6.7 Judging and synthesis (public eval)

### 6.7.1 Judge-based selection (excluding synthesis)

Judging (excluding synthesis) had a net uplift of **7** instances that were solved relative to a majority-vote baseline (selecting the most common candidate output). All 7 are minority recoveries — cases where the correct answer was not the majority output and the holistic judge identified it by reasoning over the full traces. One example of this is `dfadab01:1` (<https://arcprize.org/play?task=dfadab01>). This problem heavily suffers from “group think”: 12 of the candidate solvers converge on the same incorrect output, and another 8 converge on a second incorrect output. Only one candidate solver produces the correct output. The judges identify the originality in this lone solution and select it. The judge rationale emphasizes that most candidates recover the same underlying stamp mechanic, and that the remaining uncertainty is restricted to edge handling:

Most candidate solvers correctly identify the core stamp mechanic:

...

I prefer solutions that do \*not\* stamp at (9,8) (e.g., solutions 10-15/17-20) over those that stamp every 8.

The only remaining ambiguity is edge handling (not clearly disambiguated by the training set):

- Some solvers assume a stamp must fit fully (ignore row 17 markers).
- Others assume stamps are clipped at the border (row 17 markers produce the top 3 rows of the tile).

So the two most plausible outputs  
(same mechanic, differing only in border handling)

The judge correctly identifies that the core mechanic is not in dispute — all candidates agree on the stamp logic — and reasons that the real ambiguity lies in edge handling, which the training examples do not disambiguate. Rather than committing both guesses to the majority interpretation, the judge uses the pass@2 format to hedge: it selects one output for each interpretation of the border behavior, ensuring that the correct answer is captured regardless of which edge-handling rule the ground truth expects. This is precisely the kind of reasoning that majority voting cannot perform — it requires understanding *where* the uncertainty lies and deliberately including a minority hypothesis to cover it.

### 6.7.2 Judge synthesis

The judge is permitted to output a **synthesized** solution that is not identical to any single candidate output (Section 5.3). On the public evaluation run summarized in Figure 4, synthesis was invoked for **17** instances. Relative to an ablation that disables synthesis (i.e., forces selection among non-synthesized candidates), synthesis yields an uplift of **1** additional solved instance.

The single instance solved via synthesis is 21897d95:2 (<https://arcprize.org/play?task=21897d95>) — a particularly difficult task where **none** of the 29 candidates produced a correct output for this test instance. The task requires identifying “rooms” in the grid separated by boundaries, recognizing T-shaped arrows that encode directional recoloring instructions, and applying these instructions across room boundaries. The judge synthesized a correct output by recombining partial insights from multiple flawed candidates — identifying which candidates correctly parsed the room structure, which correctly interpreted the arrow semantics, and which correctly handled the recoloring — even though no single candidate got the full chain right. Below is an excerpt from the judge’s synthesis rationale:

```
--- COUNCIL OF JUDGES SYNTHESIZED SOLUTION ---
Across the training examples, the only solutions that
consistently match the *actual* structure are SOLUTION 1
and SOLUTION 2 (they are the same). They correctly
identify the key mechanic:
```

- The input is divided into large "rooms" (rectangular regions of a dominant background color).
  - Inside rooms, color 1 forms a 4-cell T-shape arrow (a 3-long bar + 1-cell stem).
  - The stem direction indicates which neighboring room is the target.
  - The paint color is:
    - the center of the 3-long bar if it is not 1 (a payload color), otherwise
    - the background color of the room containing the arrow.
- ...

the same room recoloring but then rotated 90 CCW.  
 This is much less likely given the square training examples, but it matches the extra rotation behavior seen in the non-square examples and is the most plausible "geometry variant" if a solver applied that step unconditionally

...

This example illustrates the core value proposition of synthesis: on tasks where the reasoning chain is too long for any single candidate to get right, the judge can act as a meta-reasoner that assembles the correct solution from complementary fragments. The current implementation remains under-optimized; further work on synthesis gating (Section 9) could increase the realized uplift, particularly on harder tasks where this recombination mechanism has the most potential.

## 6.8 Failure analysis (public eval)

Of the 39 unsolved test instances, 22 are **generation failures** (no candidate in the pool is correct) and 17 are **selection failures** (at least one correct candidate exists but is not selected by the judge). This section lists both groups to support qualitative analysis of what makes these instances hard.

### 6.8.1 Generation failures (22 instances)

The following instances have zero correct candidates across all modalities. These represent tasks where the system's hypothesis space — across text, image, and code — does not contain the ground-truth transformation.

Many of these tasks appear to require **long chains of dependent reasoning steps**, where the correct solution emerges only after composing multiple sub-concepts in sequence. Each step narrows the interpretation space, but models tend to collapse or shortcut the chain rather than faithfully executing all steps. Three illustrative examples:

- **3dc255db** (<https://arcprize.org/play?task=3dc255db>) requires recognizing “spaceships” with distinct inner and outer regions and a directional front; the outer areas are preserved while inner areas lengthen the front, with the inner/outer distinction determined by the longest edge. This involves at least four dependent inferences (shape identification, region segmentation, directional semantics, and edge-length-based classification).

- **88e364bc** (<https://arcprize.org/play?task=88e364bc>) requires interpreting a legend that encodes movement directions, then simulating movement that respects borders and colors, across multiple independently placed legend entries and areas. The combination of legend parsing, spatial simulation, and multi-region generalization defeats all modalities.
- **d35bdbdc** (test 2, <https://arcprize.org/play?task=d35bdbdc>) involves identifying a path structure, distinguishing inner from outer colors, deleting non-path elements, and recursively transferring colors — with shapes that vary in form across instances.

These examples suggest that the system’s candidate generators can sometimes identify individual sub-concepts but struggle to compose the full chain correctly. This is consistent with known limitations of LLM reasoning on deeply compositional tasks, and points to a potential benefit of multi-step decomposition approaches — though Section 8 documents that naive decomposition strategies reduced diversity in practice.

Task	Test	Link
21897d95	1	<a href="https://arcprize.org/play?task=21897d95">https://arcprize.org/play?task=21897d95</a>
2b83f449	1	<a href="https://arcprize.org/play?task=2b83f449">https://arcprize.org/play?task=2b83f449</a>
3a25b0d8	1	<a href="https://arcprize.org/play?task=3a25b0d8">https://arcprize.org/play?task=3a25b0d8</a>
3dc255db	1	<a href="https://arcprize.org/play?task=3dc255db">https://arcprize.org/play?task=3dc255db</a>
4c416de3	1	<a href="https://arcprize.org/play?task=4c416de3">https://arcprize.org/play?task=4c416de3</a>
4e34c42c	2	<a href="https://arcprize.org/play?task=4e34c42c">https://arcprize.org/play?task=4e34c42c</a>
5545f144	1	<a href="https://arcprize.org/play?task=5545f144">https://arcprize.org/play?task=5545f144</a>
6ffbe589	1	<a href="https://arcprize.org/play?task=6ffbe589">https://arcprize.org/play?task=6ffbe589</a>
88e364bc	1	<a href="https://arcprize.org/play?task=88e364bc">https://arcprize.org/play?task=88e364bc</a>
8b7bacbf	1	<a href="https://arcprize.org/play?task=8b7bacbf">https://arcprize.org/play?task=8b7bacbf</a>
8b7bacbf	2	<a href="https://arcprize.org/play?task=8b7bacbf">https://arcprize.org/play?task=8b7bacbf</a>
9bbf930d	1	<a href="https://arcprize.org/play?task=9bbf930d">https://arcprize.org/play?task=9bbf930d</a>
a25697e4	1	<a href="https://arcprize.org/play?task=a25697e4">https://arcprize.org/play?task=a25697e4</a>
abc82100	1	<a href="https://arcprize.org/play?task=abc82100">https://arcprize.org/play?task=abc82100</a>
b9e38dc0	1	<a href="https://arcprize.org/play?task=b9e38dc0">https://arcprize.org/play?task=b9e38dc0</a>
d35bdbdc	2	<a href="https://arcprize.org/play?task=d35bdbdc">https://arcprize.org/play?task=d35bdbdc</a>
da515329	1	<a href="https://arcprize.org/play?task=da515329">https://arcprize.org/play?task=da515329</a>
dbff022c	1	<a href="https://arcprize.org/play?task=dbff022c">https://arcprize.org/play?task=dbff022c</a>
de809cff	1	<a href="https://arcprize.org/play?task=de809cff">https://arcprize.org/play?task=de809cff</a>
e12f9a14	1	<a href="https://arcprize.org/play?task=e12f9a14">https://arcprize.org/play?task=e12f9a14</a>
e12f9a14	2	<a href="https://arcprize.org/play?task=e12f9a14">https://arcprize.org/play?task=e12f9a14</a>
faa9f03d	1	<a href="https://arcprize.org/play?task=faa9f03d">https://arcprize.org/play?task=faa9f03d</a>

### 6.8.2 Selection failures (17 instances)

The following instances have at least one correct candidate, but the holistic judge fails to select it. These represent cases where the judge’s selection mechanism leads to an incorrect final output.

A qualitative observation across several selection failures is that the **test instance introduces a new mechanic or configuration not fully disambiguated by the training examples**, creating genuine ambiguity about the correct generalization. In these cases, some candidates make the right assumptions about how the rule extends, but the judge — reasoning from the same ambiguous training pairs — tends to favor the majority interpretation. Two illustrative examples:

**36a08778** (test 2, <https://arcprize.org/play?task=36a08778>): The training examples establish a straightforward “water flows downward” mechanic. Test 2 introduces walls that block flow — a new structural element not present in training. Most candidates (and the judge) converge on the simpler mechanic, discarding the minority candidates that correctly handle walls:

```
Most of the 29 solvers converged to the *same* mechanic (and the same output)
...
I discarded outliers like **Solution 2** and **Solution 1**
```

**88bcf3b4** (test 2, <https://arcprize.org/play?task=88bcf3b4>): Training examples show a single “rope/snake” component moving in one direction. Test 2 introduces multiple strings moving in multiple directions — a generalization the training pairs do not disambiguate. The judge identifies the ambiguity but cannot resolve it, and selects the wrong resolution:

```
From the 5 training examples, the consistent mechanic is **not** "gravity" or
"attraction" of whole blobs. Instead, one non-background component acts like a
**rope/snake**
```

```
...
The main ambiguity the examples do *not* disambiguate is what happens when the
"returning" segment reaches the pole's column/row again:
```

- **Solution 20** continues the diagonal return even after hitting the pole's column (so it can pass "through" and go beyond).
- **Solution 16** effectively **clamps** once aligned with the pole's column (keeps going straight instead of overshooting).

This failure mode is structurally difficult for the holistic judge: when training examples are consistent with multiple generalizations and the test instance is the only signal that distinguishes them, the judge faces the same underspecification that makes the task hard in the first place. The judge’s tendency to favor majority-cluster agreement — which is beneficial on tasks where the majority is correct (e.g., `dfadab01:1` in Section 6.7) — becomes a liability on tasks where the correct generalization is a minority hypothesis precisely because it requires handling a novel test-time mechanic.

Task	Test	Link
16b78196	1	<a href="https://arcprize.org/play?task=16b78196">https://arcprize.org/play?task=16b78196</a>
35ab12c3	1	<a href="https://arcprize.org/play?task=35ab12c3">https://arcprize.org/play?task=35ab12c3</a>
36a08778	2	<a href="https://arcprize.org/play?task=36a08778">https://arcprize.org/play?task=36a08778</a>
4c7dc4dd	1	<a href="https://arcprize.org/play?task=4c7dc4dd">https://arcprize.org/play?task=4c7dc4dd</a>
4e34c42c	1	<a href="https://arcprize.org/play?task=4e34c42c">https://arcprize.org/play?task=4e34c42c</a>
6e4f6532	1	<a href="https://arcprize.org/play?task=6e4f6532">https://arcprize.org/play?task=6e4f6532</a>
7666fa5d	1	<a href="https://arcprize.org/play?task=7666fa5d">https://arcprize.org/play?task=7666fa5d</a>
78332cb0	1	<a href="https://arcprize.org/play?task=78332cb0">https://arcprize.org/play?task=78332cb0</a>
78332cb0	2	<a href="https://arcprize.org/play?task=78332cb0">https://arcprize.org/play?task=78332cb0</a>
7b80bb43	1	<a href="https://arcprize.org/play?task=7b80bb43">https://arcprize.org/play?task=7b80bb43</a>
88bcf3b4	2	<a href="https://arcprize.org/play?task=88bcf3b4">https://arcprize.org/play?task=88bcf3b4</a>
89565ca0	1	<a href="https://arcprize.org/play?task=89565ca0">https://arcprize.org/play?task=89565ca0</a>
9aaea919	1	<a href="https://arcprize.org/play?task=9aaea919">https://arcprize.org/play?task=9aaea919</a>
a32d8b75	1	<a href="https://arcprize.org/play?task=a32d8b75">https://arcprize.org/play?task=a32d8b75</a>

Task	Test	Link
d35bdbdc	1	<a href="https://arcprize.org/play?task=d35bdbdc">https://arcprize.org/play?task=d35bdbdc</a>
e3721c99	2	<a href="https://arcprize.org/play?task=e3721c99">https://arcprize.org/play?task=e3721c99</a>
eee78d87	1	<a href="https://arcprize.org/play?task=eee78d87">https://arcprize.org/play?task=eee78d87</a>

## 7 Ablation Studies

A rigorous component attribution would require multiple full-pipeline runs with individual components removed or substituted. Each such run costs approximately \$2,400 in API spend (Table 5), making extensive ablation prohibitively expensive — a full ablation matrix across modalities, judge configurations, and candidate budgets would cost tens of thousands of dollars. This paper therefore relies primarily on **post-hoc analysis of the single public evaluation run**, extracting what can be measured from the existing data (e.g., comparing judge selections against majority-vote baselines on the same candidate pool) rather than running dedicated ablation experiments. This approach has clear limitations: it cannot capture interaction effects between components, and some comparisons (e.g., end-to-end modality removal) require fresh runs that have not been performed. The ablations reported here should be read with this constraint in mind.

Unless otherwise stated, “solved” refers to pass@2 at the **test-instance** level on the public evaluation split (167 instances).

### 7.1 Measured ablations

#### 7.1.1 Judging: holistic selection vs per-candidate scoring (excluding synthesis)

Judging (excluding synthesis) had a net uplift of **7** solved instances relative to a **majority-vote baseline** that selects the most common candidate output grid as the first guess and the second-most-common as the second guess (i.e., standard self-consistency). All 7 uplift instances are **minority recoveries**: cases where the correct answer was not the most frequent candidate output, and the holistic judge identified it by reasoning over the full traces rather than counting votes. This directly validates the “anti-consistency” motivation (Section 5): on these tasks, the majority cluster was wrong, and the holistic judge’s ability to read and compare reasoning traces — rather than simply tallying outputs — was the deciding factor. One illustrative instance is **dfadab01:1** (<https://arcprize.org/play?task=dfadab01>), where the candidate pool clusters around two distinct incorrect hypotheses and the holistic judge selects the lone correct candidate.

#### 7.1.2 Judging: synthesis enabled vs disabled

Synthesized solution yields a net uplift of **1** additional solved instance in this run. Across all three judges, synthesis was invoked **17** times total; in most cases the synthesized output did not change the final selected solution, as the weighted scoring still favored non-synthesized candidates.

While the measured uplift is small on this particular run, synthesis has been a more material contributor in other experimental runs during development. The mechanism has particular potential on harder tasks where no single candidate is fully correct but multiple candidates contain complementary partial insights — exactly the regime where recombination should help most. Better understanding when and how to trigger synthesis (see “synthesis gating” in Section 9) is an important direction for future work.

**Table 4. Measured judge ablations on the public evaluation run.** Reported deltas are net solved-instance uplifts attributable to the indicated component (with the stated control condition).

Component	Ablation / control condition	Net uplift (solved instances)
Holistic selection	Holistic judge vs majority-vote baseline (synthesis disabled in both settings)	+7 (all minority recoveries)
Judge synthesis	Synthesis enabled vs disabled (holistic selection held fixed)	+1

### 7.1.3 Cost attribution per component

Table 5 reports the cost breakdown for the public evaluation run, averaged per test instance (i.e., per task-test-pair, of which there are 167). This is distinct from the per-task cost reported in Table 1 (\$19.69/task over 120 tasks), because many tasks contain multiple test instances (75 tasks have 1 test instance, 43 have 2, and 2 have 3). The per-instance average (\$14.31) does not sum to the per-task figure because the per-task metric aggregates all test instances within a task into a single cost. Note: a strict roll-up of the Table 5 total (\$2,390.28 / 120 tasks = \$19.92/task) differs slightly from the reported \$19.69/task; the discrepancy is likely due to accumulated floating-point rounding in the cost-accounting script.

**Table 5. Cost attribution per test instance on the public evaluation run (n = 167).**

Phase	Total (\$)	Avg \$/instance	% of total
Candidate generation	2081.37	12.46	87.1%
Judging	308.91	1.85	12.9%
<b>Total</b>	<b>2390.28</b>	<b>14.31</b>	<b>100%</b>

**Table 6. Candidate generation cost by modality family.**

Modality family	Total (\$)	Avg \$/instance	% of generation cost
Text (incl. deep think)	597.70	3.58	28.7%
Image	467.10	2.80	22.5%
Code	1016.56	6.09	48.9%

Candidate generation dominates overall cost at 87% of spend, with judging accounting for only 13%. Within generation, code candidates are the most expensive family (49% of generation cost), driven primarily by tool-integrated code generation with iterative sandbox execution. Given that judging contributes +7 solved instances (holistic selection) and +1 (synthesis) at only 13% of total cost, the judging phase is highly cost-effective relative to its accuracy contribution.

## 7.2 Modality ablations (oracle-level only)

Section 6.5 reports modality-level uniqueness on the **complete-coverage subset** (n = 130), where all modalities are executed. In this subset, the following **exclusive** oracle solvability counts are observed (Table 3): Text only = 2, Image only = 6, Code only = 7. These exclusive counts imply

that removing any single modality would reduce candidate-oracle coverage by at least a few percent even before accounting for downstream judge interactions and selection effects.

However, oracle-level analysis has an important limitation: it measures whether a correct candidate *exists* in a modality’s output, but does not measure the end-to-end effect of removing that modality on the final system output. Removing a modality could affect judge behavior in ways not captured by oracle overlap — for instance, reducing the number of candidates changes cluster dynamics, which could make it easier or harder for the holistic judge to identify the correct solution. A modality might also contribute “near-miss” candidates that inform judge synthesis even when no individual candidate from that modality is exactly correct.

The proper ablation — running the full pipeline with one modality family removed and re-running judging on the reduced candidate pool — has not been performed. Each such run requires regenerating all judge transcripts on the reduced candidate set (and ideally multiple repetitions to account for judge variance), making it expensive relative to the oracle-level analysis. This remains a gap; the oracle-level uniqueness numbers in Section 6.5 should be interpreted as a lower bound on each modality’s contribution, not as a precise end-to-end attribution.

### 7.3 Unperformed ablations

The following ablations would strengthen the paper’s claims but have not been run due to cost constraints. They are listed here both as transparency about what remains unknown and as a roadmap for future work.

#### Generation ablations:

- **End-to-end modality removal:** run the full pipeline with one modality family (text, image, or code) removed and re-run judging on the reduced candidate pool. The oracle-level exclusive counts in Section 7.3 provide a lower bound, but the actual end-to-end impact — including judge interaction effects — is unknown. This requires at minimum three full runs ( $\sim \$7,200$  total).
- **Independent candidates vs sequential refinement:** hold compute fixed and compare  $N$  independent candidates (the current approach) against  $N$  sequential refinement steps (iterative prompt chaining or staged decomposition). Section 8 documents qualitative evidence that sequential approaches reduced diversity, but a controlled comparison with matched compute budgets has not been performed.
- **Candidate budget scaling:** sweep the number of candidates per modality/model to estimate marginal returns per additional candidate and identify diminishing-returns regimes. The current system uses 29 candidates, but it is unknown whether 15 or 50 would yield meaningfully different accuracy at proportionally different cost.
- **Per-model contribution:** isolate the contribution of each foundation model (GPT-5.2, Gemini 3, Opus 4.5) by running the pipeline with one model removed entirely. Opus 4.5 contributes only 1 candidate; whether its inclusion is cost-justified relative to adding another GPT-5.2 or Gemini candidate is unknown.
- **Temperature and sampling parameters:** the current system uses default or near-default sampling settings for each model. Sweeping temperature, top-p, and other sampling parameters within each modality could reveal whether diversity is better increased through sampling variation or through modality variation.
- **Representation formats:** CSV vs alternative encodings (e.g., JSON-like arrays, Python list syntax, and object-abstraction encodings), evaluated under the same candidate/judge budgets. The benchmarking reported in Section 4 was performed during development with

different model versions; a controlled evaluation on the final system would be more rigorous.

#### Selection and judging ablations:

- **Full majority-vote baseline comparison:** the +7 holistic selection uplift reported above is computed post hoc by comparing judge selections against the majority output grid on the same candidate pool. A cleaner comparison would run majority vote as the *sole* selection mechanism in a full end-to-end run (without any judge invocation), eliminating any confounds from shared infrastructure. This would also quantify whether the 13% cost of judging is justified by the accuracy gain.
- **Judge ensemble sizing:** the current system uses 3 judges with weighted scoring. No ablation comparing 1-judge vs 3-judge accuracy has been performed. Quantifying the disagreement rate and how often the ensemble corrects vs overrides individual judges would clarify whether the ensemble cost is justified or whether a single judge suffices.
- **Alternative selection mechanisms:** compare the holistic judge against other selection strategies on the same candidate pool, including per-output log-probability scoring, pairwise judge tournaments (comparing candidates two at a time), and best-of-N with a reward model. These comparisons would isolate the contribution of joint-context evaluation from other factors.
- **Judge model diversity:** the current system uses three GPT-5.2 judges. A mixed-model ensemble (e.g., one GPT-5.2, one Gemini, one Opus judge) was tested informally and underperformed (Section 5), but a rigorous comparison — controlling for prompt format and scoring calibration — has not been done.
- **Trace content ablation:** the holistic judge receives full reasoning traces alongside candidate outputs. Comparing judge accuracy with traces vs outputs-only would quantify whether the trace content actually helps selection or whether the judge primarily relies on output grid comparison.

#### Early stopping ablations:

- **Early stopping threshold tuning:** the current heuristic triggers early stopping when initial candidates agree. Varying the agreement threshold and the number of candidates consulted before the stopping decision would characterize the accuracy/cost trade-off and the groupthink risk documented in Section 6.6.
- 

## 8 Negative Results and Discarded Approaches

This section documents explored approaches that were ultimately discarded, often because they reduced diversity, increased brittleness, or forced premature abstraction.

### 8.1 Hint generation → solver (discarded)

This approach is structurally similar to iterative self-improvement methods such as Self-Refine (Madaan et al. 2023) and Reflexion (Shinn et al. 2023), where an initial pass generates feedback that informs a subsequent attempt.

Motivation: - doubling the reasoning budget across two turns (hint then solve).

Observed drawback: - the hint stage often **limits creativity** and collapses candidate diversity into a narrower space, which is counterproductive when trying to break new ground.

## 8.2 Object identification → transformation identification → solver (discarded)

Motivation: - structured decomposition to “force” abstraction.

Failure mode: - brittle handoff between stages. Both verbose and overly terse handovers caused confusion and reduced diversity, often regressing toward the mean rather than expanding the hypothesis space.

## 8.3 Opus codegen and Opus image reasoning (discarded from final mix)

In the final system, Opus contributes only a single text-reasoning candidate. Opus codegen and image reasoning were tested but contributed less uniquely relative to the GPT/Gemini configurations, and were dropped from the final candidate mix.

## 8.4 Grid representations and output constraints (discarded variants)

Key findings:

- **CSV-style** encoding outperformed many alternatives, especially as grids grow and representation consumes a large share of the reasoning budget.
- Forcing strict outputs (e.g., **requiring JSON** via API-level response formats) underperformed. This is important because strict schemas are a common LLM engineering practice, but appear to reduce model effectiveness on this domain.

Engineering trade-off: - removing constraints increases output noise; robust parsing (regex + validation) becomes necessary, but was worth it for accuracy.

## 8.5 Synthetic data augmentation for code candidates (discarded)

Motivation: - generate additional training examples (e.g., via color permutation, rotation, or mirroring of the provided pairs) to give code-generation candidates more test cases to validate against, potentially improving program correctness.

Reasons for discarding: - **Surface-level augmentations add little signal.** Color permutations produce nominally distinct examples but do not test new structural properties of the transformation; a program that overfits to the original examples will typically also pass color-permuted variants. - **Geometric transforms break semantics.** Rotation and mirroring alter the meaning of tasks that depend on absolute orientation — for example, gravity-based tasks (where “down” matters) or tasks where spatial relationships across training pairs encode the rule (e.g., map reconstruction). Applying these transforms indiscriminately would introduce incorrect training signal. - **Meaningful augmentation requires solving the task first.** Generating genuinely informative synthetic examples (new inputs paired with correct outputs) requires knowing the transformation rule — which is the problem the solver is trying to solve. This makes meaningful augmentation infeasible in a private-dataset evaluation setting where ground-truth transformations are unavailable.

## 8.6 Extensive prompt engineering (discarded)

This is perhaps the most counterintuitive finding in the paper, and it directly contradicts standard LLM engineering practice.

The conventional approach to LLM integration treats the model as a programmable API: the more precisely you specify the desired behavior — step-by-step instructions, prescribed reasoning

templates, output schemas, chain-of-thought scaffolding — the better the results. This works well for structured tasks like data extraction, classification, or format conversion, where the solution space is well-defined and the model’s job is to comply with a specification.

On ARC-AGI, this approach consistently degraded performance. During development, I tested numerous prompt engineering strategies including:

- **Prescribed reasoning templates:** instructing the model to first identify objects, then describe transformations, then apply them step by step.
- **Structured output requirements:** requiring specific output formats (e.g., JSON grids via API-level response format constraints).
- **Detailed chain-of-thought scaffolding:** breaking the reasoning into named stages (“Step 1: Identify the pattern. Step 2: Describe the rule. Step 3: Apply to test input.”).
- **Domain-specific heuristics in the prompt:** suggesting the model look for symmetry, rotation, color mapping, or other common ARC patterns.
- **Iterative prompt refinement:** tuning prompt wording based on failure analysis of specific tasks.

In every case, the more prescriptive the prompt, the worse the system performed on the hardest tasks. The final system uses a deliberately minimal prompt (Section 4) that provides only the task data, a brief context sentence, and a request to explain reasoning — with no prescribed structure, no step-by-step template, and no domain heuristics.

The mechanism appears to be a **compliance tax on reasoning**: when the model is given detailed instructions about *how* to think, it allocates a significant portion of its reasoning budget to following those instructions rather than to actually solving the problem. On easy tasks — where the transformation is simple enough that any reasonable approach works — this overhead is harmless. On hard tasks — where the model needs to make creative leaps, entertain unusual hypotheses, or reason about structures it has never seen — the overhead is fatal. The model dutifully follows the prescribed template, produces a well-formatted but wrong answer, and never explores the unstructured reasoning path that might have led to the correct solution.

This also interacts with diversity: a prescriptive prompt narrows the hypothesis space across candidates. When all N candidates follow the same reasoning template, they tend to converge on the same (possibly wrong) answer. A minimal prompt allows different candidates to approach the problem from genuinely different angles, increasing the probability that at least one candidate finds the correct transformation.

The implication is that for novel reasoning tasks — where the solution is not known in advance and the model must discover it — **the best prompt is often the least prescriptive one**. The engineer’s job shifts from programming the model’s behavior to removing obstacles to the model’s reasoning. This is uncomfortable for systems engineers accustomed to treating prompts as specifications, but on this benchmark, letting the model think freely and tolerating noisy outputs (with robust downstream parsing) consistently outperformed carefully engineered prompts.

---

## 9 Discussion and Conclusion

This paper demonstrates that strong ARC-AGI-2 performance can be achieved by treating **modalities as search operators** and pairing diverse candidate generation with context-preserving selection:

- generate candidates independently across heterogeneous reasoning channels (text, image, code) to maximize hypothesis diversity,
- then select using holistic judging over full traces.

## 9.1 Limitations

This work has several important limitations that should be weighed when interpreting the results.

**Cost and practical scalability.** The system spends \$19.69 per task on the public evaluation and \$38.99 per task on the semi-private set — orders of magnitude more expensive than a single model call. ARC Prize explicitly treats efficiency as a first-class metric (Chollet et al. 2024), and the current system is far from efficient. The brute-force strategy of generating 29 candidates across three models and three modalities, then running three separate judge passes over the full candidate pool, is effective but wasteful. Much of the cost is spent on candidates that contribute nothing to the final answer. A production system would need adaptive routing — spending heavily only on tasks that resist cheap methods — but no such routing mechanism has been developed here.

**Single-run results.** The headline numbers (76.11% public, 72.9% semi-private) each come from a single evaluation run. LLM outputs are stochastic, and the system’s reliance on sampling diversity means that results will vary across runs. No confidence intervals are reported because repeated full-pipeline runs were not performed (each costing ~\$2,400). The true expected accuracy could be meaningfully higher or lower than the reported figures, and the variance is unknown.

**Incomplete ablation coverage.** Section 7 reports post-hoc ablations for holistic selection (+7 instances) and synthesis (+1 instance), but these are extracted from a single run rather than from controlled experiments with matched baselines. Several ablations that would substantially strengthen the paper’s claims — end-to-end modality removal, independent vs sequential generation, candidate budget scaling, judge ensemble sizing — have not been performed due to cost constraints. The component attribution claims in this paper should be understood as indicative rather than rigorous. Section 7 provides a detailed list of unperformed ablations.

**Reproducibility fragility.** The system depends on specific proprietary model snapshots (GPT-5.2, Gemini 3 Preview, Opus 4.5) and vendor-specific “reasoning settings” (e.g., OpenAI’s “x-high” reasoning effort) that may change or become unavailable over time. Model providers routinely update model weights, deprecate API parameters, and alter rate limits without notice. Results may not be exactly reproducible even with identical prompts and parameters if the underlying model versions drift. To maximize reproducibility within these constraints, the full source code is open-sourced at <https://github.com/beetree/ARC-AGI> and the complete raw data for the public evaluation run — including all prompts, responses, reasoning traces, and judge transcripts (over 7 million lines) — is available on Kaggle at <https://www.kaggle.com/code/johanland/johanland-solver-v7-public/comments?scriptVersionId=290052212>. Exact replication is nonetheless not guaranteed.

**No learning across tasks.** The system treats each task independently — no information is carried from one task to the next. A human solver would build intuitions across tasks (e.g., “tasks in this benchmark often involve symmetry” or “I’ve seen this color-mapping pattern before”), but the current system starts from scratch every time. This is both a limitation and a design choice: task independence simplifies the pipeline and avoids overfitting to task ordering, but it means the system cannot amortize its reasoning cost across related tasks.

**Narrow evaluation domain.** The results are demonstrated on a single benchmark (ARC-AGI-2). While the architectural pattern — diverse generation plus holistic judging — is domain-general in principle, this paper provides no evidence that the approach transfers to other domains. The specific design choices (modality mix, candidate count, judge prompt structure) were tuned

for ARC and may not generalize without adaptation.

## 9.2 Future work

- Adaptive routing: allocate expensive modalities only when uncertainty is high.
- Judge compression without premature abstraction: find ways to reduce context size while retaining the benefits of joint context.
- Further ablations: Section 7 lists a detailed set of unperformed ablations — including judge ensemble sizing, trace content contribution, per-model attribution, and early-stopping threshold tuning — that would strengthen component attribution claims.
- Synthesis gating and amplification: the current judge always has the option to synthesize a novel output. A gating mechanism that decides *when* synthesis is likely to help — e.g., only when candidate agreement is low or when no candidate passes a confidence threshold — could improve targeting. Conversely, when synthesis is identified as having high potential (e.g., multiple candidates contain complementary partial solutions), the system could invoke additional synthesis attempts with varied prompting to increase the probability of a correct recombination. The current single-pass synthesis yielded +1 instance (Section 7); a more aggressive, targeted synthesis strategy could yield further uplift on the hardest tasks where no single candidate is fully correct.
- Image representation tuning: the finding that intentionally imprecise grid renderings outperform pixel-perfect ones (Section 4) is suggestive but not well understood. Systematic study of rendering parameters — resolution, distortion level, color palette, annotation style — and their interaction with different vision-language models could yield further gains and clarify when and why visual prompting helps.
- Broader modality coverage: additional frontier providers and open-source models, plus parameter sweeps (temperature, etc.).
- Formal diversity quantification: the current paper measures modality complementarity via oracle overlap counts (Tables 2–3), but a richer diversity measure — e.g., pairwise output disagreement rates, embedding-space distances between reasoning traces, or information-theoretic metrics over the candidate distribution — would enable principled decisions about which generators to add, remove, or scale. A task-archetype taxonomy (classifying ARC tasks by the type of reasoning they require) could further clarify which modalities are most valuable for which problem classes.
- Domain transfer: the “diverse generation + holistic judging” pattern is not specific to ARC. Any domain where models produce confident but divergent answers — mathematical proof search, legal analysis, medical diagnosis — could benefit from context-preserving adjudication over multiple independent reasoning traces. Validating this on non-ARC benchmarks is a natural next step.

## 9.3 A note on AI-assisted development

The solver was designed and implemented entirely through AI tools (Section 3.4), with the author acting as strategic director rather than programmer. This creates a recursive dynamic: frontier LLMs were used to design, debate, and implement a system whose purpose is to push the frontier of what LLMs can solve. That this approach produced a competitive result — without the author writing or reading any solver code — suggests that AI-mediated systems engineering is becoming a viable development paradigm for complex, novel pipelines.

## 9.4 Reproducibility

The full source code for the solver is available at: <https://github.com/beetree/ARC-AGI>. The repository contains all prompts, tool schemas, candidate generation configurations, and judging logic.

The complete public-evaluation run — including all API parameters, model versions, and raw logs (prompts, responses, reasoning traces, intermediate artifacts, and judge transcripts; over 7 million lines) — is available as a Kaggle notebook: <https://www.kaggle.com/code/johanland/johanland-solver-v7-public/comments?scriptVersionId=290052212>.

The semi-private evaluation was executed by ARC Prize’s verification infrastructure; the author does not control that environment and cannot release those logs.

## 9.5 Acknowledgments

Thank you to the ARC-AGI Discord community for valuable discussion and shared insights throughout the development of this work, and to Greg Kamradt at ARC Prize for conducting the official semi-private evaluation.

## 9.6 Closing

ARC-AGI-2 progress is moving quickly, and the benchmark is explicitly designed to push beyond what scaling alone yields. The approach here aims to contribute a practical, reproducible systems pattern for advancing that frontier:

### Search across modalities, judge in full context.<sup>8</sup>

- Akyürek, Ekin, Mehul Damani, Adam Zweiger, Linlu Qiu, Han Guo, Jyothish Pari, Yoon Kim, and Jacob Andreas. 2024. “The Surprising Effectiveness of Test-Time Training for Few-Shot Learning.” *arXiv Preprint arXiv:2411.07279*.
- Besta, Maciej, Nils Blach, Ales Kubicek, Robert Gerstenberger, Michal Podstawski, Lukas Giani-nazzi, Joanna Gajber, et al. 2024. “Graph of Thoughts: Solving Elaborate Problems with Large Language Models.” In *Proceedings of the AAAI Conference on Artificial Intelligence*. Vol. 38.
- Chollet, François. 2019. “On the Measure of Intelligence.” *arXiv Preprint arXiv:1911.01547*.
- Chollet, François, Mike Knoop, Gregory Kamradt, and Bryan Landers. 2024. “ARC Prize 2024: Technical Report.” *arXiv Preprint arXiv:2412.04604*.
- Ferré, Sébastien. 2021. “First Steps of an Approach to the ARC Challenge Based on Descriptive Grid Models and the Minimum Description Length Principle.” *arXiv Preprint arXiv:2112.00848*.
- Fletcher-Hill, Paul. 2024. “Mini-ARC: Solving Abstraction and Reasoning Puzzles with Small Transformer Models.”
- Gao, Luyu, Aman Madaan, Shuyan Zhou, Uri Alon, Pengfei Liu, Yiming Yang, Jamie Callan, and Graham Neubig. 2023. “PAL: Program-Aided Language Models.” *International Conference on Machine Learning*.
- Hodel, Michael. 2024. “Addressing the Abstraction and Reasoning Corpus via Procedural Example Generation.” *arXiv Preprint arXiv:2404.07353*.
- Li, Wen-Ding, Keya Hu, Carter Larsen, Yuqing Wu, Simon Alford, Caleb Woo, Spencer M. Dunn, et al. 2024. “Combining Induction and Transduction for Abstract Reasoning.” *arXiv Preprint arXiv:2411.02272*.

<sup>8</sup>This paper was drafted with Codex CLI and completed with Claude Code, under the supervision of Johan Land. Peer review was conducted by a parallel Claude Code instance, ChatGPT-5.2-Pro, and Gemini 3 Deep Think. Any remaining errors are, naturally, the human’s fault.

- Macfarlane, Matthew V., and Clément Bonnet. 2024. “Searching Latent Program Spaces.” *arXiv Preprint arXiv:2411.08706*.
- Madaan, Aman, Niket Tandon, Prakhar Gupta, Skyler Hallinan, Luyu Gao, Sarah Wiegreffe, Uri Alon, et al. 2023. “Self-Refine: Iterative Refinement with Self-Feedback.” *Advances in Neural Information Processing Systems* 36.
- Moffitt, Michael D. 2025. “ARC-GEN: A Mimetic Procedural Benchmark Generator for the Abstraction and Reasoning Corpus.” *arXiv Preprint arXiv:2511.00162*.
- Moskvichev, Arseny, Victor Vikram Odouard, and Melanie Mitchell. 2023. “The ConceptARC Benchmark: Evaluating Understanding and Generalization in the ARC Domain.” *Transactions on Machine Learning Research*.
- Ouellette, Simon. 2024. “Towards Efficient Neurally-Guided Program Induction for ARC-AGI.” *arXiv Preprint arXiv:2411.17708*.
- Puget, Jean-François. 2024. “A 2D nGPT Model for ARC Prize.”
- Schick, Timo, Jane Dwivedi-Yu, Roberto Dessì, Roberta Raileanu, Maria Lomeli, Luke Zettlemoyer, Nicola Cancedda, and Thomas Scialom. 2023. “Toolformer: Language Models Can Teach Themselves to Use Tools.” In *Advances in Neural Information Processing Systems*. Vol. 36.
- Shinn, Noah, Federico Cassano, Ashwin Gopinath, Karthik Narasimhan, and Shunyu Yao. 2023. “Reflexion: Language Agents with Verbal Reinforcement Learning.” *Advances in Neural Information Processing Systems* 36.
- Snell, Charlie, Jaehoon Lee, Kelvin Xu, and Aviral Kumar. 2024. “Scaling LLM Test-Time Compute Optimally Can Be More Effective Than Scaling Model Parameters.” *arXiv Preprint arXiv:2408.03314*.
- Wang, Xuezhi, Jason Wei, Dale Schuurmans, Quoc Le, Ed Chi, Sharan Narang, Aakanksha Chowdhery, and Denny Zhou. 2023. “Self-Consistency Improves Chain of Thought Reasoning in Language Models.” In *International Conference on Learning Representations*.
- Wei, Jason, Xuezhi Wang, Dale Schuurmans, Maarten Bosma, Brian Ichter, Fei Xia, Ed Chi, Quoc Le, and Denny Zhou. 2022. “Chain-of-Thought Prompting Elicits Reasoning in Large Language Models.” *Advances in Neural Information Processing Systems* 35.
- Yao, Shunyu, Dian Yu, Jeffrey Zhao, Izhak Shafran, Thomas L. Griffiths, Yuan Cao, and Karthik Narasimhan. 2023. “Tree of Thoughts: Deliberate Problem Solving with Large Language Models.” In *Advances in Neural Information Processing Systems*. Vol. 36.
- Yao, Shunyu, Jeffrey Zhao, Dian Yu, Nan Du, Izhak Shafran, Karthik Narasimhan, and Yuan Cao. 2023. “ReAct: Synergizing Reasoning and Acting in Language Models.” *International Conference on Learning Representations*.
- Zheng, Lianmin, Wei-Lin Chiang, Ying Sheng, Siyuan Zhuang, Zhanghao Wu, Yonghao Zhuang, Zi Lin, et al. 2023. “Judging LLM-as-a-Judge with MT-Bench and Chatbot Arena.” *Advances in Neural Information Processing Systems* 36.